

# BAG OF WORDS (BOW)

Agro-IODAA, semestre 1

Vincent Guigue

# Handling textual data: the classification case

- 1 Big corpus  $\Leftrightarrow$  Huge vocabulary
- 2 Sentence structure is hard to model
- 3 Words are polymorphous: singular/plural, masculine/feminine
- 4 Synonyms: how to deal with?
- 5 Machine learning + large dimensionality = problems

# Handling textual data: the classification case

- 1 Big corpus  $\Leftrightarrow$  Huge vocabulary

**Perceptron, SVM, Naive Bayes... Boosting, Bagging...**

**Distributed & efficient algorithms**

- 2 Sentence structure is hard to model

**Removing the structure...**

- 3 Words are polymorphous: singular/plural, masculine/feminine

**Several approaches... (see below)**

- 4 Synonyms: how to deal with?

**wait for the next course !**

- 5 Machine learning + large dimensionality = problems

**Removing useless words**

# BAG OF WORDS

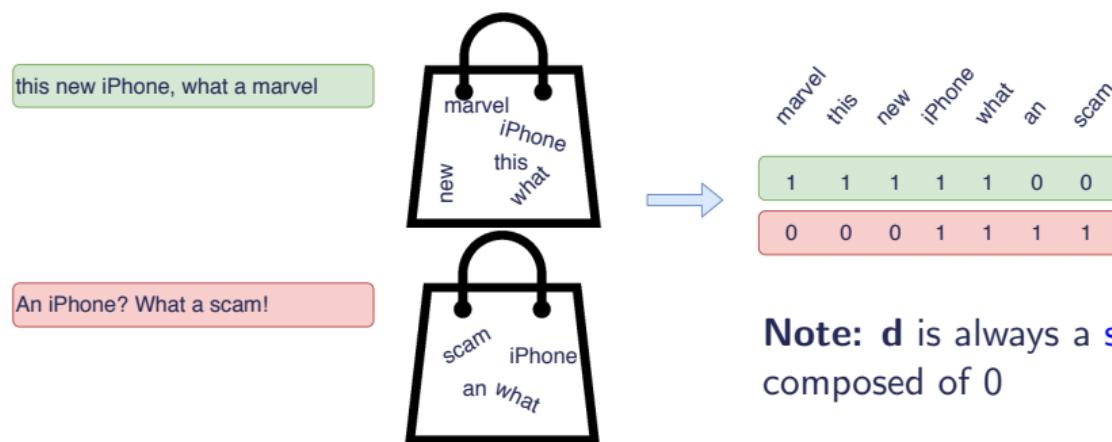
# Bag of words

Sentence structure = costly handling

⇒ Elimination !

## *Bag of words* representation

- 1 Extraction of vocabulary  $V$
- 2 Each document becomes a counting vector :  $\mathbf{d} \in \mathbb{N}^{|V|}$



**Note:**  $\mathbf{d}$  is always a **sparse vectors**, mainly composed of 0

# Example

Set of toy documents:

```
1  documents = ['The lion does not live in the jungle', \
2  'Lions eat big preys', \
3  'In the zoo, the lion sleep', \
4  'Self-driving cars will be autonomous in towns', \
5  'The future car has no steering wheel', \
6  'My car already has sensors and a camera']
```

Dictionary

The lion does not live in the jungle Lions eat big preys In zoo, sleep Self-driving cars will be autonomous towns future car has no steering wheel My already sensors and a camera

Green level  $\propto$  nb occurrences

## Information coding

## Counting words appearing in 2 documents:

The lion does not live in the jungle  
In the zoo, the lion sleep

The does in jungle lion live not the Lions big eat preys In sleep zoo, Self-driving autonomous be cars towns will car future has no steering wheel My a already and camera sensors

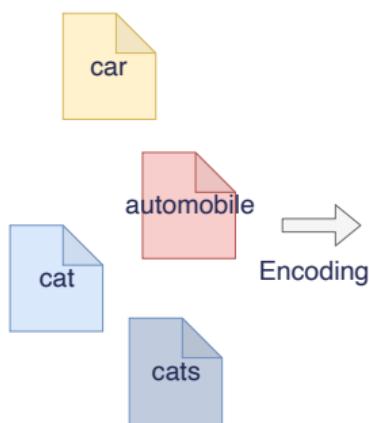
- + We are able to vectorize textual information
  - Dictionary requires preprocessing

# Word representation & semantic gap

All words are orthogonal:

Considering virtual 2 documents made of a single word :

$$\begin{bmatrix} 0 & \dots & 0 & d_{ik} > 0 & \dots & 0 \\ 0 & \dots & d_{jk'} > 0 & 0 & \dots & 0 \end{bmatrix}$$



Vectorized corpus			
d1	1	0	0
d2	0	0	1
d3	0	1	0

word 1 ... car ... automobile ... cat ... word D

Same distances

Then:  $k \neq k' \Rightarrow \mathbf{d}_i \cdot \mathbf{d}_j = 0$

...even if  $w_k = lion$  and  $w_{k'} = lions$

⇒ Definition of the **semantic gap**

No metrics between words



## Semantic issue

Understanding documents = matching relevant descriptors

- Syntactic difference ⇒ orthogonality of the representation vectors
- Word groups : more intrinsic semantics...
  - ... but fewer match with other document
  - N-grams ⇒ dictionary size ↗
  - N-grams = great potential...
    - but require careful preprocessings

*This film was not interesting*

- Unigrams: this, film, was, not, interesting
- bigrams: this\_film, film\_was, was\_not, not\_interesting
- N-grams... + combination: e.g. 1-3 grams



# Text format

## ■ Bag-of-Words, BoW

### + Advantages

- Easy, light
- fast
- Opportunity to enrich
- Efficient on document classification
- Efficient implementations

(Real-time systems, IR (indexing)...)

(POS, context encoding, N-gram...)

`nltk, sklearn`

### - Drawbacks

- Loose document/sentence structure

⇒ **Several tasks almost impossible to tackle**

- NER, POS tagging, SRL
- Text generation

# DIMENSIONALITY OF THE BAG OF WORDS



# Vocabulary size & word occurrences

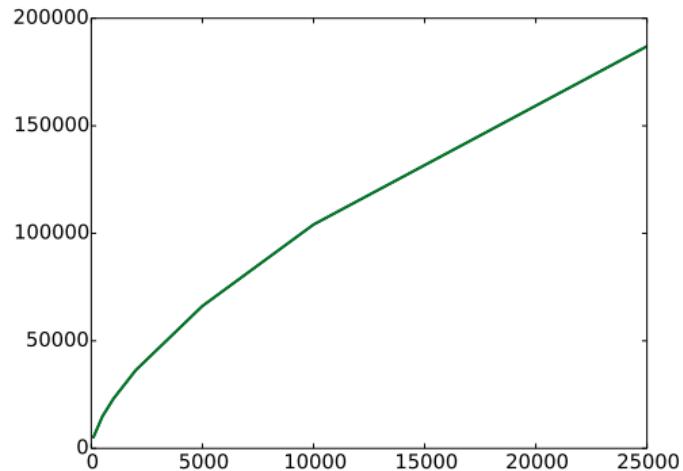
## Vocabulary size:

$$|V| \propto \log(N), \quad N = \text{number of documents}$$

On movie reviews :

$|V|$  with respect to # reviews  
Let's have a closer look on the axes !!!

25k docs  $\Leftrightarrow$  200k words !!



# Vocabulary size & word occurrences

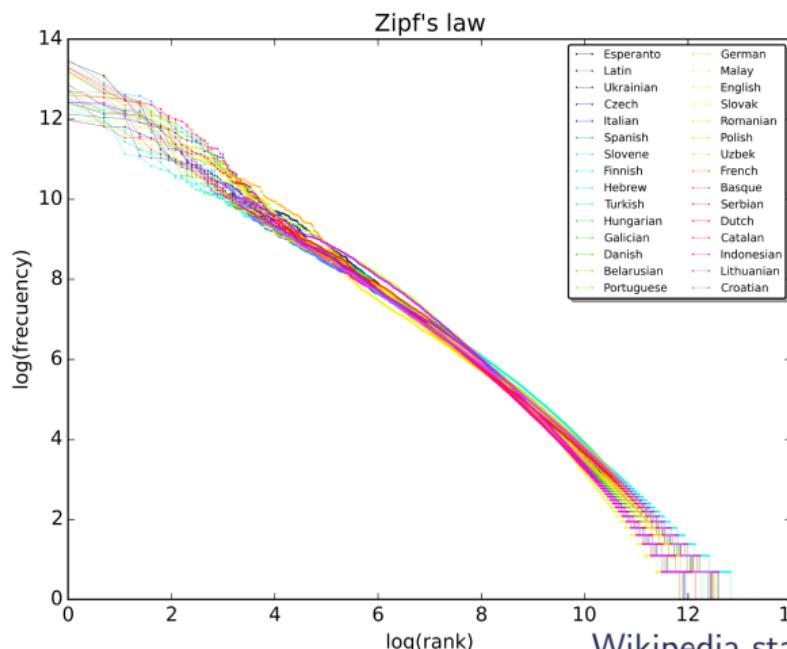
## Vocabulary size:

$$|V| \propto \log(N), \quad N = \text{number of documents}$$

## Word occurrence distribution:

*Frequency*; of word; in corpus C

*Rank*; : frequency rank ( $\approx$  popularity) of word;



# Vocabulary size & word occurrences

## Vocabulary size:

$$|V| \propto \log(N), \quad N = \text{number of documents}$$

For example, in the *Brown Corpus of American English* text, the word "**the**" is the most frequently occurring word, and by itself accounts for nearly **7%** of all word occurrences (69,971 out of slightly over 1 million). True to Zipf's law, the second-place word "**of**" accounts for slightly over **3.5%** of words (36,411 occurrences), followed by "**and**" (28,852).

[Wikipedia](#)

# Vocabulary size & word occurrences

## Vocabulary size:

$$|V| \propto \log(N), \quad N = \text{number of documents}$$

## Corpus size:

Which weight for a  $25k \times 200k$  matrix of double?

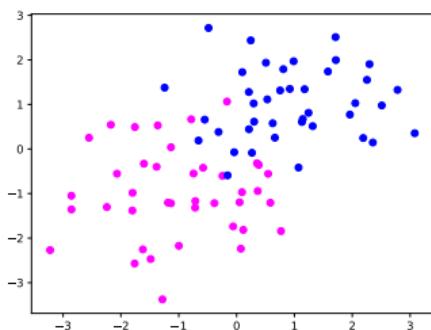
$$= 5 \cdot 10^9 \times 4\text{bytes}$$

⇒ Need a specific data structure: **Sparse Matrix**

# Curse of dimensionality

A classical toy example to illustrate the curse of dimensionality:

Original dataset:



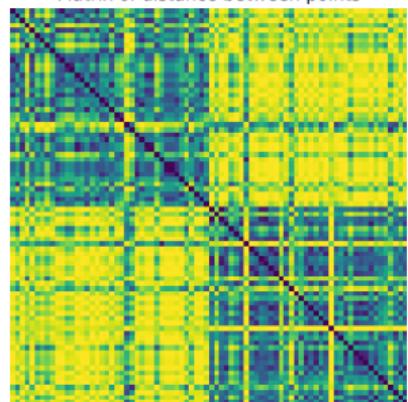
Matrix view

Matrix of raw points



Distance matrix

Matrix of distance between points

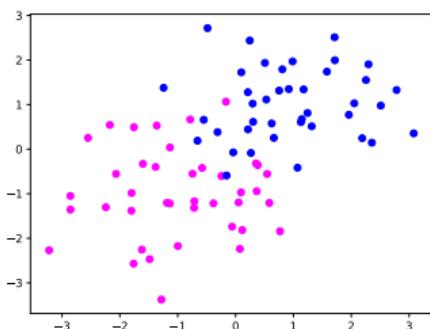


Easy problem / classes are clearly separated

# Curse of dimensionality

A classical toy example to illustrate the curse of dimensionality:

Original dataset:



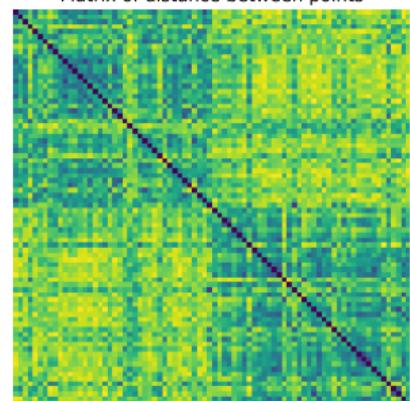
Matrix view

Matrix of raw points



Distance matrix

Matrix of distance between points

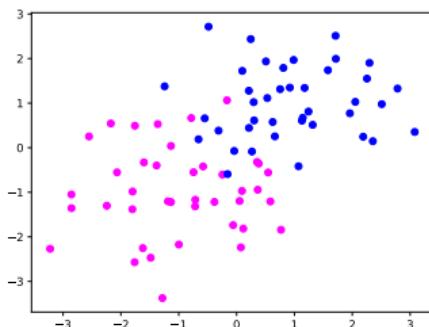


Adding some noisy dimensions in the dataset

# Curse of dimensionality

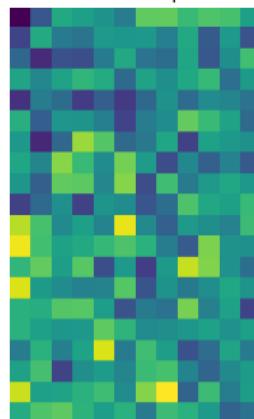
A classical toy example to illustrate the curse of dimensionality:

Original dataset:



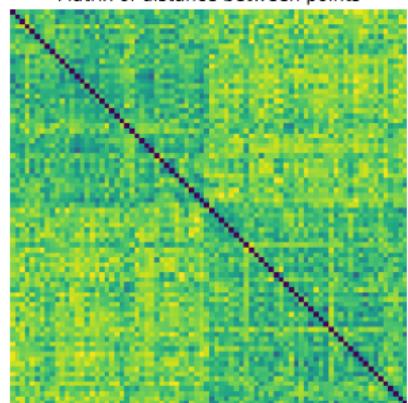
Matrix view

Matrix of raw points



Distance matrix

Matrix of distance between points



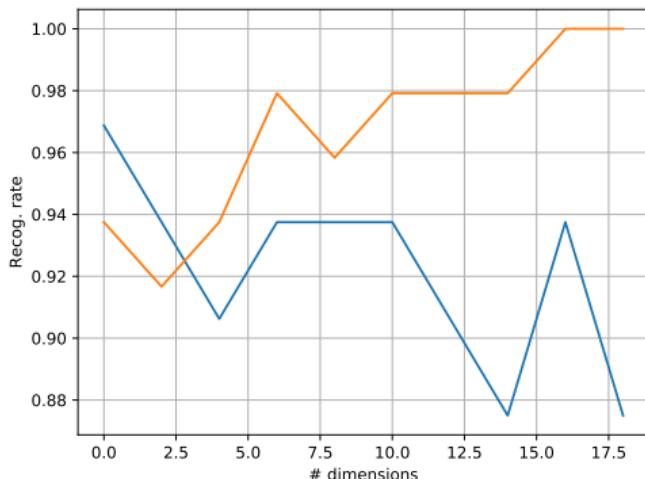
Adding **more** noisy dimensions in the dataset

⇒ Euclidian distance is very sensitive to the dimensionality issue

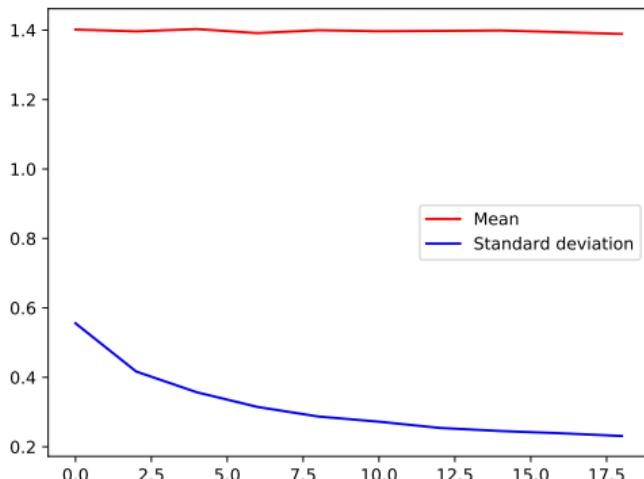
# Curse of dimensionality

A classical toy example to illustrate the curse of dimensionality:

Basic classifier on those datasets



Distances between points in the dataset



⇒ Learn accuracy ↗, test accuracy ↘ = **overfitting**

⇒ All points tend to lay on an hypersphere (they become equidistant)

# Dimensionality (and noise) reduction

```
1 documents = [ 'The lion does not live in the jungle', \
2   'Lions eat big preys', \
3   'In the zoo, the lion sleep', \
4   'Self-driving cars will be autonomous in towns', \
5   'The future car has no steering wheel', \
6   'My car already has sensors and a camera']
```

*Original dictionary:*



# Dimensionality (and noise) reduction

```
1 documents = [ 'The lion does not live in the jungle', \
2   'Lions eat big preys', \
3   'In the zoo, the lion sleep', \
4   'Self-driving cars will be autonomous in towns', \
5   'The future car has no steering wheel', \
6   'My car already has sensors and a camera']
```

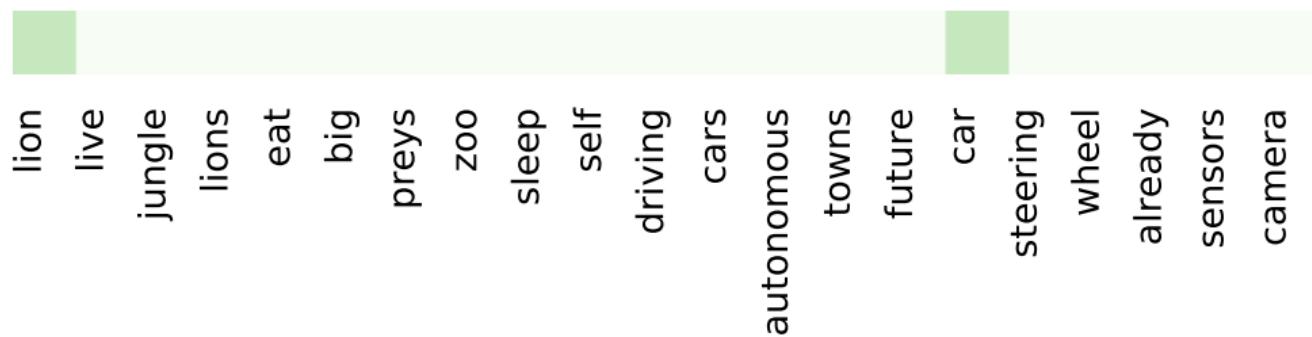
*Removing capitals:*

the lion does not live in jungle lions eat big preys zoo, sleep self-driving cars will be autonomous towns future car has no steering wheel my already sensors and a camera

# Dimensionality (and noise) reduction

```
1 documents = ['The lion does not live in the jungle', \
2 'Lions eat big preys', \
3 'In the zoo, the lion sleep', \
4 'Self-driving cars will be autonomous in towns', \
5 'The future car has no steering wheel', \
6 'My car already has sensors and a camera']
```

*Removing stop words:*



*Implementation:* black list (nltk) or upper frequency bound

# Dimensionality (and noise) reduction

---

```
1 documents = [ 'The lion does not live in the jungle', \
2   'Lions eat big preys', \
3   'In the zoo, the lion sleep', \
4   'Self-driving cars will be autonomous in towns', \
5   'The future car has no steering wheel', \
6   'My car already has sensors and a camera' ]
```

---

*Removing rare words* occurring less than a threshold :

Dictionary = {lion, car}

⇒ too extreme in this toy example...

But a good idea in real situations.

Remainder: rare words represent the a large part of the dictionary

⇒ Tricky setting of the thresholds (upper & lower bounds)

# Dimensionality (and noise) reduction

```
1 documents = [ 'The lion does not live in the jungle', \
2   'Lions eat big preys', \
3   'In the zoo, the lion sleep', \
4   'Self-driving cars will be autonomous in towns', \
5   'The future car has no steering wheel', \
6   'My car already has sensors and a camera']
```

## Lemmatization:

in linguistics is the process of grouping together the **inflected forms** of a word so they can be analysed as a single item, identified by the **word's lemma**, or dictionary form

⇒ Requires advanced linguistic resources including words and inflected forms.  
E.g. : lions ⇒ lion; are ⇒ be; ...

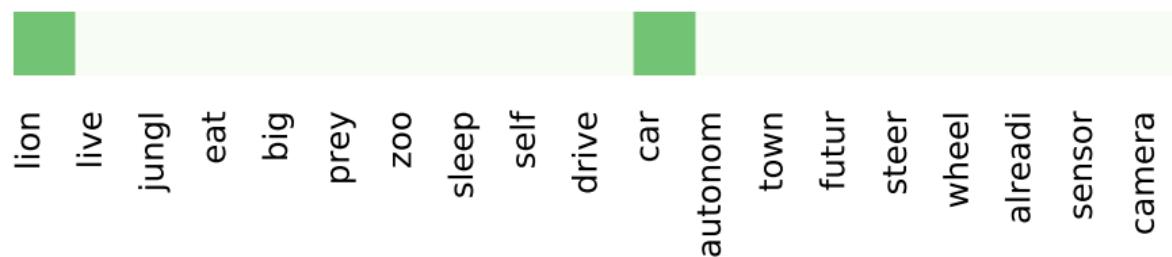
# Dimensionality (and noise) reduction

```
1 documents = ['The lion does not live in the jungle', \
2   'Lions eat big preys', \
3   'In the zoo, the lion sleep', \
4   'Self-driving cars will be autonomous in towns', \
5   'The future car has no steering wheel', \
6   'My car already has sensors and a camera']
```

## Stemming:

A statistical approximated process of the lemmatization

E.g. : removing s or ly at the end of the words

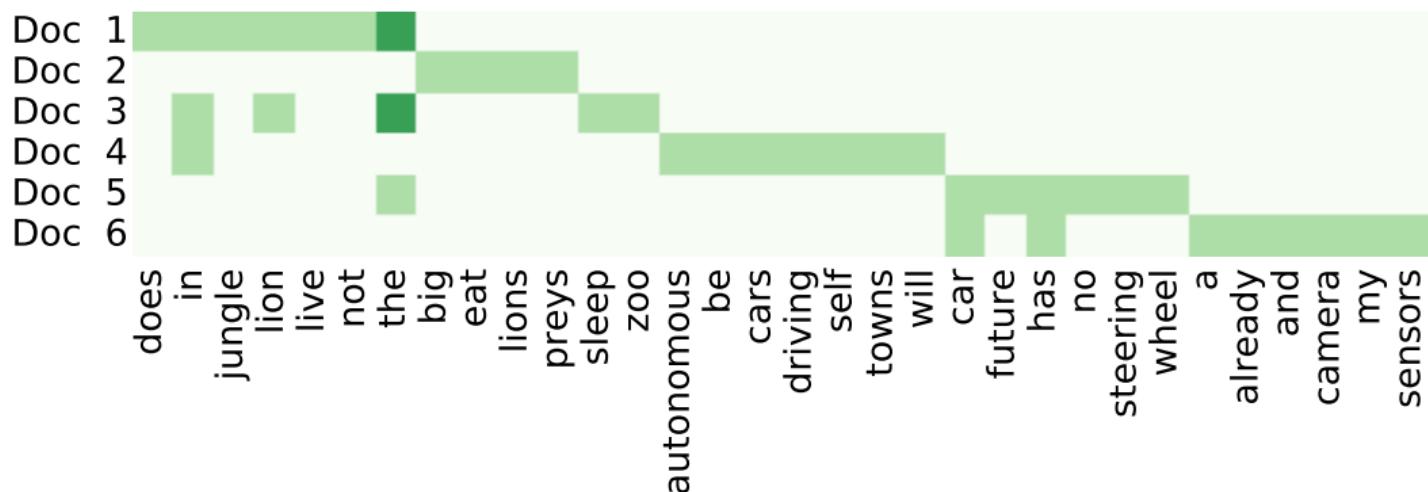


**Note:** it is not a problem to create invalid words... If they are stable!

# Corpus representation

```
1 documents = ['The lion does not live in the jungle', \
2 'Lions eat big preys', \
3 'In the zoo, the lion sleep', \
4 'Self-driving cars will be autonomous in towns', \
5 'The future car has no steering wheel', \
6 'My car already has sensors and a camera']
```

Corpus mapping on the basic dictionary:



## Metrics between documents

Given documents  $\mathbf{d}_i \in \mathbb{R}^{|D|}$  and  $\mathbf{d}_j \in \mathbb{R}^{|D|}$  with the dictionary  $D$ .

First idea : **Euclidian metrics**

$$d(\mathbf{d}_i, \mathbf{d}_j) = \|\mathbf{d}_i - \mathbf{d}_j\| = \sqrt{\sum_k (d_{ik} - d_{jk})^2}$$

But:

$$d(\mathbf{d}_i, \mathbf{d}_j) = \sqrt{\|\mathbf{d}_i\|^2 + \|\mathbf{d}_j\|^2 - 2\mathbf{d}_i \cdot \mathbf{d}_j}$$

⇒ Sensitive to the norm of  $\mathbf{d}$  or to the ratio  $\mathbf{d}_i \cdot \mathbf{d}_j$  vs  $\|\mathbf{d}\|$

# Metrics between documents

- **Euclidian distance**  
⇒ not robust enough
- **Inner product**

$$\text{sim}(\mathbf{d}_i, \mathbf{d}_j) = \frac{\mathbf{d}_i \cdot \mathbf{d}_j}{\|\mathbf{d}_i\| \|\mathbf{d}_j\|} = \cos(\widehat{\mathbf{d}_i}, \widehat{\mathbf{d}_j}) \propto \sum_k d_{ik} d_{jk}$$

⇒ focusing on common non-zeros dimensions

- **Kullback Leibler**

Assuming that each document can be seen as a distribution over words (ie,  
 $\forall i, \sum_k d_{ik} = 1$ )

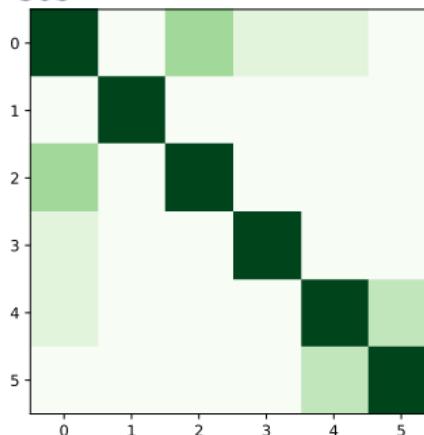
$$D_{\text{KL}}(\mathbf{d}_i \parallel \mathbf{d}_j) = \sum_k d_{ik} \log \frac{d_{ik}}{d_{jk}}$$

⇒ not very stable, take care of  $d_{ik} = 0$  or  $d_{jk} = 0$

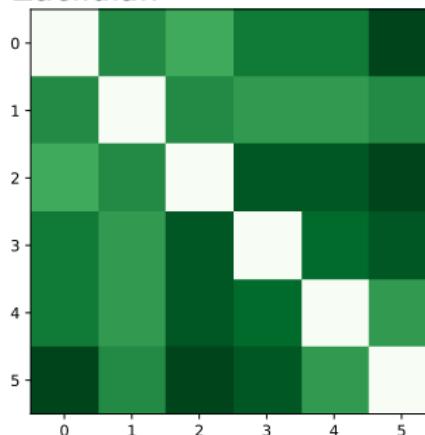
# A Metrics on our example

```
1 documents = [ 'The lion does not live in the jungle', \
2   'Lions eat big preys', \
3   'In the zoo, the lion sleep', \
4   'Self-driving cars will be autonomous in towns', \
5   'The future car has no steering wheel', \
6   'My car already has sensors and a camera']
```

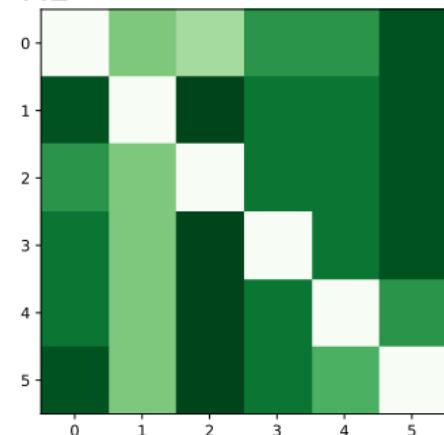
Cos



Euclidian



KL

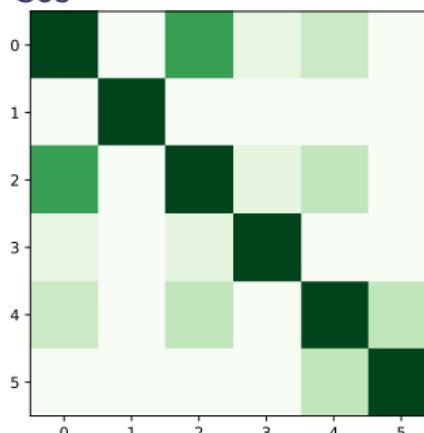


Basic representation of texts... **Too much noise !**

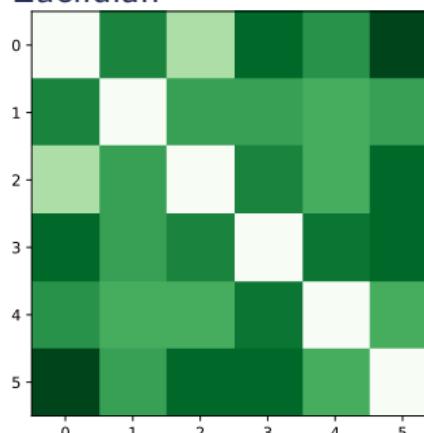
# A Metrics on our example

```
1 documents = ['The lion does not live in the jungle', \
2 'Lions eat big preys', \
3 'In the zoo, the lion sleep', \
4 'Self-driving cars will be autonomous in towns', \
5 'The future car has no steering wheel', \
6 'My car already has sensors and a camera']
```

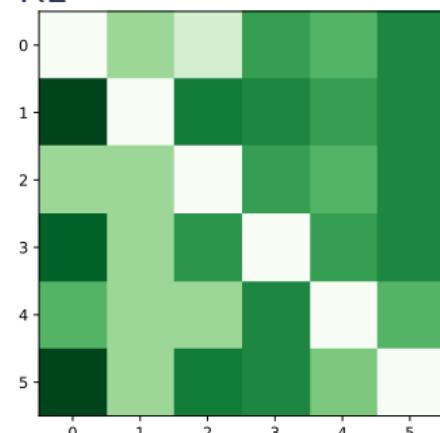
Cos



Euclidian



KL

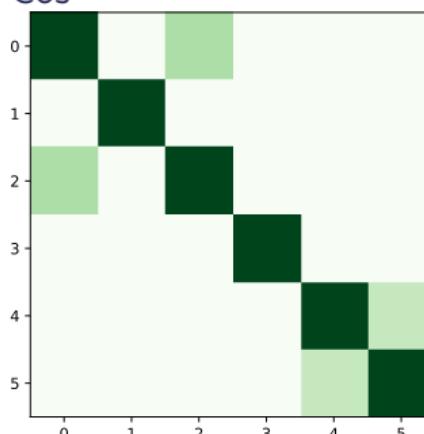


Preprocessing (removing capitals/punctuation)  $\Rightarrow$  situation still confuse

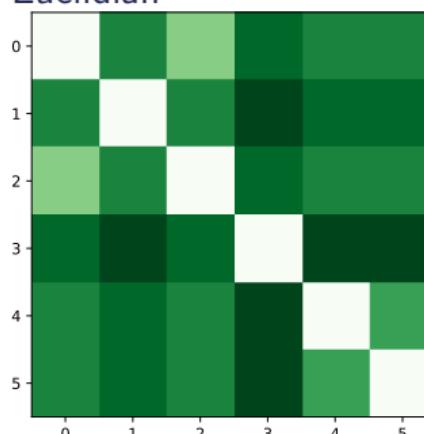
# A Metrics on our example

```
1 documents = ['The lion does not live in the jungle', \
2 'Lions eat big preys', \
3 'In the zoo, the lion sleep', \
4 'Self-driving cars will be autonomous in towns', \
5 'The future car has no steering wheel', \
6 'My car already has sensors and a camera']
```

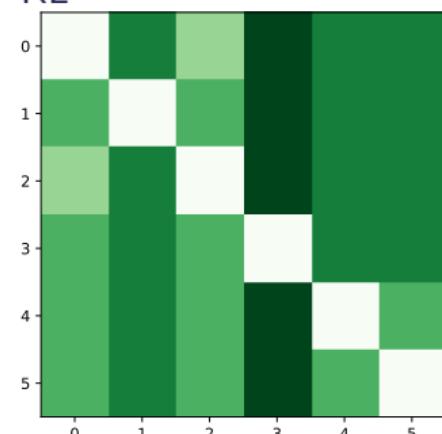
Cos



Euclidian



KL

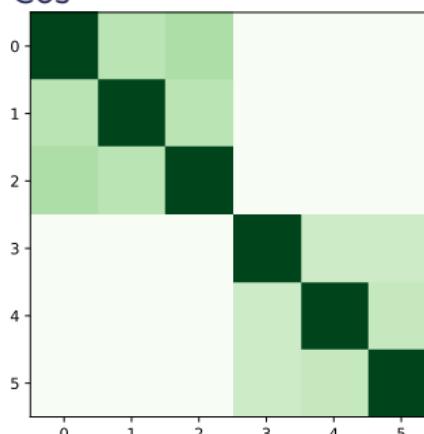


Preprocessing + removing stop words ⇒ slight improvement

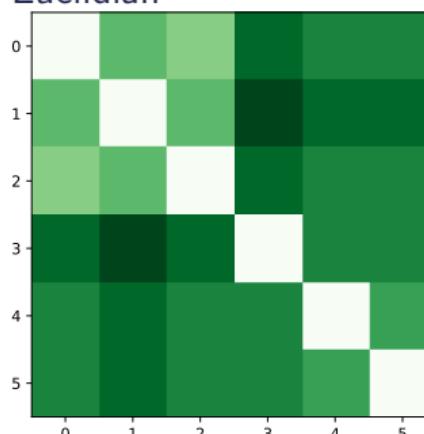
# A Metrics on our example

```
1 documents = ['The lion does not live in the jungle', \
2 'Lions eat big preys', \
3 'In the zoo, the lion sleep', \
4 'Self-driving cars will be autonomous in towns', \
5 'The future car has no steering wheel', \
6 'My car already has sensors and a camera']
```

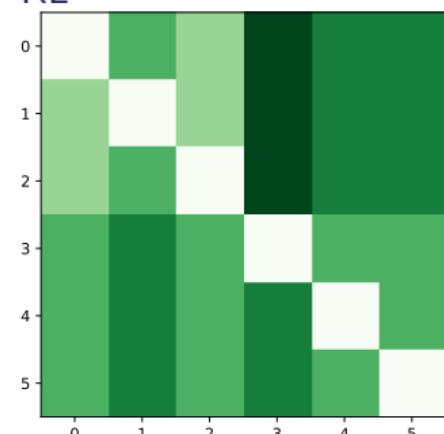
Cos



Euclidian



KL



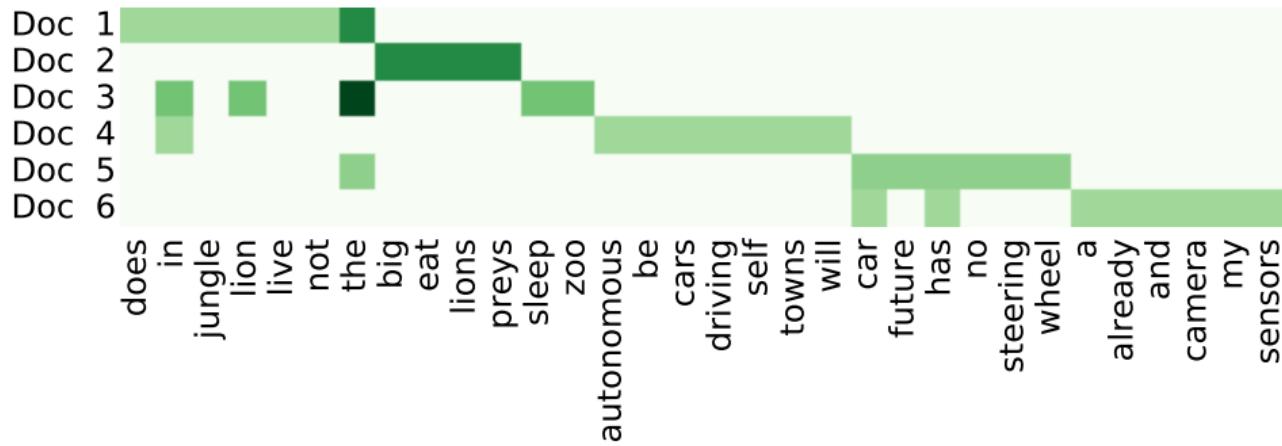
Preprocessing + removing stop words + stemming ⇒

distinction between classes

# Classical issues in text modeling

Every document in the corpus has the **same nearest neighbor**...

A strong magnet with many words (=long document)



**Normalization**  $\Rightarrow$  descriptors = **term frequency** in the document

$$\forall i, \quad \sum_k d_{ik}^{(tf)} = 1$$

# Classical issues in text modeling

Frequent words are overweighted...

*if word k is in most documents, it is probably useless*

Introduction of the **document frequency**:

$$df_k = \frac{|\{\mathbf{d} : t_k \in \mathbf{d}\}|}{|C|}, \quad \text{Corpus : } C = \{\mathbf{d}_1, \dots, \mathbf{d}_{|C|}\}$$

Tf-idf coding = term frequency, inverse document frequency:

$$d_{ik}^{(tfidf)} = d_{ik}^{(tf)} \log \frac{|C|}{|\{\mathbf{d} : t_k \in \mathbf{d}\}|}$$

⇒ A strong idea to focus on **keywords**...

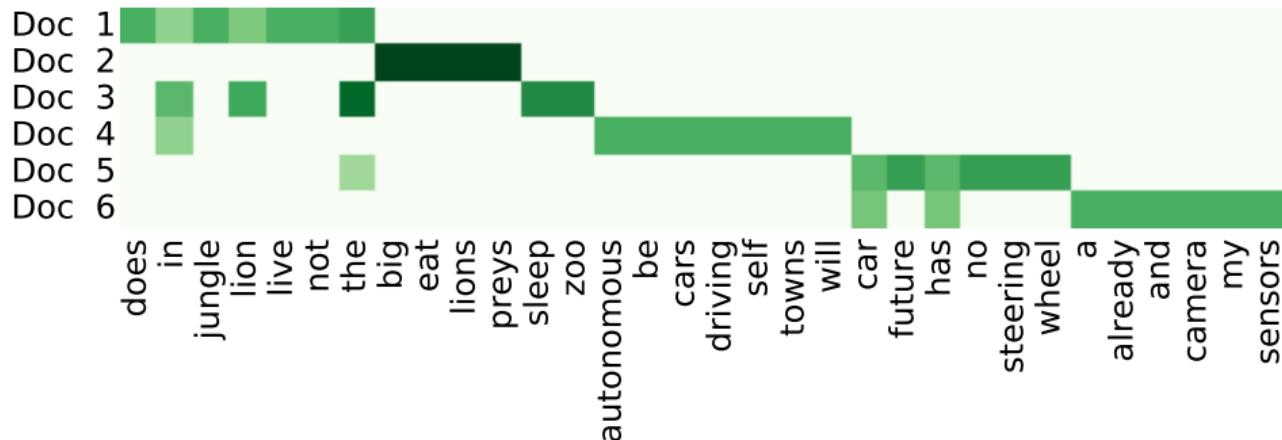
... But not strong enough to get rid of all stop words

⇒ TFIDF should be combined with blacklists

## Classical issues in text modeling

Frequent word are overweighted...

*if word k is in most documents, it is not discriminant*

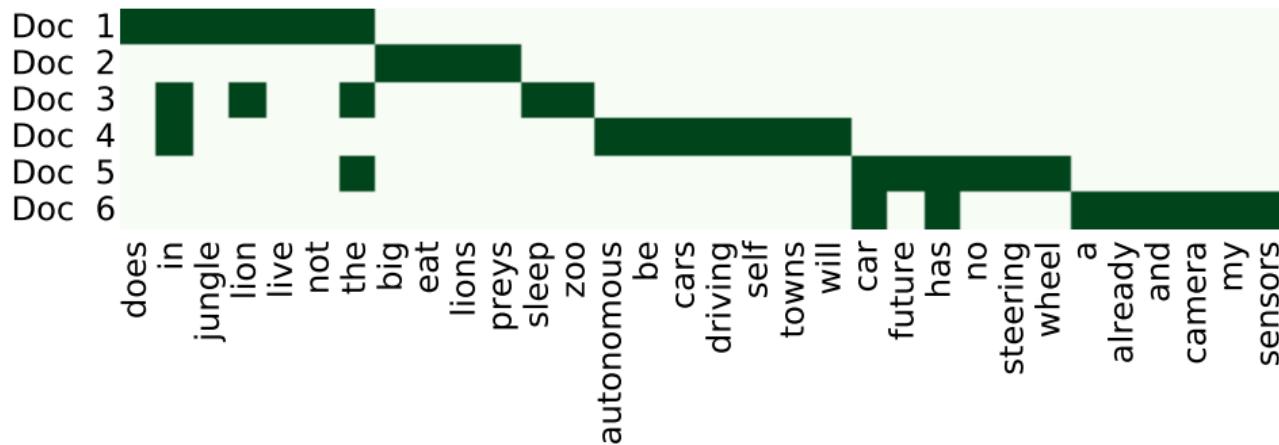


$$d_{ik}^{(tfidf)} = d_{ik}^{(tf)} \log \frac{|C|}{|\{\mathbf{d} : t_k \in \mathbf{d}\}|}$$

# Classical issues in text modeling

In some specific cases (e.g. sentiment classification), it is more robust to remove all information about frequency...

⇒ Presence coding



$$d_{ik}^{(pres)} = \begin{cases} 1 & \text{if word } k \text{ is in } d_i \\ 0 & \text{else} \end{cases}$$



# A small digression onto Information Retrieval

IR main task :

Answering a query  $\mathbf{q} = \{q_1, \dots, q_n\}$  by selecting documents  $\mathbf{d}$  according to metrics :  
 $dist(\mathbf{q}, \mathbf{d})$

Most common metrics: BM25

$$\text{score}(\mathbf{q}, \mathbf{d}) = \sum_{i=1}^n \text{IDF}(q_i) \cdot \frac{\text{freq}(q_i, \mathbf{d}) \cdot (k_1 + 1)}{\text{freq}(q_i, \mathbf{d}) + k_1 \cdot \left(1 - b + b \cdot \frac{|\mathbf{d}|}{\text{avgdl}}\right)}$$

$$\text{IDF}(q_i) = \log \frac{N - n(q_i) + 0.5}{n(q_i) + 0.5}, \quad b = 0.75, k_1 \in [1.2, 2.0]$$

$N$  : Corpus size,       $n(q_i)$  : Number of documents with  $q_i$

IR subtasks:

- Enforcing serendipity
- Modeling source authority (PageRank)

# Dimensionality reduction as a learning problem

- Eliminating word according to a criterion (still preprocessing)
  - Saliency :  $S_{tf-idf}(i) = \frac{\sum_j tf\text{-}idf(i,j)}{|\{tf\text{-}idf(i,j) \neq 0\}|}$  (word  $i$ , word  $j$ )
  - Odds ratio:  $S_{odds}(i) = \frac{p_i/(1-p_i)}{q_i/(1-q_i)} = \frac{p_i(1-q_i)}{q_i(1-p_i)}$ . (often in log). Where  $p_i$  is the frequency of  $t_i$  in class 1 and  $q_i$  is the frequency of  $t_i$  in class 2.
  - Other criteria :Fisher, Mallows... Based on **separability**
- Regularization (improving robustness in learning)
  - cf after

# DOCUMENT CLASSIFICATION

## Avant les traitements, l'encodage...

<http://sametmax.com/lencoding-en-python-une-bonne-fois-pour-toute/>

- Sur le disque, les fichiers sont encodés de manière spécifique...
- En python 2, les strings sont encodées de manière spécifique...

## Avant les traitements, l'encodage...

<http://sametmax.com/lencoding-en-python-une-bonne-fois-pour-toute/>

- Sur le disque, les fichiers sont encodés de manière spécifique...
- En python 2, les strings sont encodées de manière spécifique...
- L'ouverture des fichiers est souvent associée à un encodage !!!

⇒ Comment gérer cela?

# Avant les traitements, l'encodage...

<http://sametmax.com/lencoding-en-python-une-bonne-fois-pour-toute/>

- Sur le disque, les fichiers sont encodés de manière spécifique...
- En python 2, les strings sont encodées de manière spécifique...
- L'ouverture des fichiers est souvent associée à un encodage !!!

⇒ Comment gérer cela?

## Solution 1

- 1 Ouverture en binaire des fichiers (e.g. en python)
- 2 Conversion des strings depuis un encodage connu  
`str.decode('utf8')`  
`unicodedata, unidecode`

## Solution 2

- 1 Vérifier le type d'encodage + convertir avant

# Chaine de traitements standard

## 1. Preprocessing

- encodage (latin, utf8, ...)
- ponctuation
- stemming
- lemmatisation
- tokenization
- minuscule/maj
- regex
- ...

## 2. Mise en forme

- Construction d'un dictionnaire (index)
- + Index inversé (pour l'explication des traitements)
- **Mise en forme vectorielle**
- **Conservation des séquences**

## 3. Traitements

- Classification des docs, des mots, des phrases
- Sémantique
- ...  
Perceptron ou HMM?

# Exemples de tâches

## ■ Classification **thématique**

- classer les news sur un portail d'information,
- trier des documents pour la veille sur internet,
- présenter les résultats d'une requête

## ■ Classification **d'auteurs**

- review spam,
- détection d'auteurs

## ■ Information pertinente/non pertinente

- filtrage personnalisé (à partir d'exemple), classifieur actif (évoluant au fil du temps)
- spam/non spam

## ■ Classification de **sentiments**

- documents positifs/négatifs, sondages en ligne

# Naive Bayes

- Très rapide, interprétable:  
le classifieur *historique* pour les sacs de mots

# Naive Bayes

- Très rapide, interprétable:  
le classifieur *historique* pour les sacs de mots
- Modèle génératif:
  - ensemble des documents  $\{d_i\}_{i=1,\dots,N}$ ,
  - documents = une suite de mots  $w_j$ :  $d_i = (w_1, \dots, w_{|d_i|})$ .
  - modèle  $\Theta_c$  pour chaque classe de documents.
  - max de vraisemblance pour l'affectation

# Naive Bayes

- Très rapide, interprétable:  
le classifieur *historique* pour les sacs de mots
- Modèle génératif:
  - ensemble des documents  $\{d_i\}_{i=1,\dots,N}$ ,
  - documents = une suite de mots  $w_j$ :  $d_i = (w_1, \dots, w_{|d_i|})$ .
  - modèle  $\Theta_c$  pour chaque classe de documents.
  - max de vraisemblance pour l'affectation
- Modélisation naïve:  $P(d_i|\Theta_c) = \prod_{j=1}^{|d_i|} P(w_j|\Theta_c) = \prod_{j=1}^{|D|} P(w_j|\Theta_c)^{x_i^j}$   
 $x_i^j$  décrit le nombre d'apparitions du mot  $j$  dans le document  $i$

# Naive Bayes

- Très rapide, interprétable:  
le classifieur *historique* pour les sacs de mots
- Modèle génératif:
  - ensemble des documents  $\{d_i\}_{i=1,\dots,N}$ ,
  - documents = une suite de mots  $w_j$ :  $d_i = (w_1, \dots, w_{|d_i|})$ .
  - modèle  $\Theta_c$  pour chaque classe de documents.
  - max de vraisemblance pour l'affectation
- Modélisation naïve:  $P(d_i|\Theta_c) = \prod_{j=1}^{|d_i|} P(w_j|\Theta_c) = \prod_{j=1}^{|D|} P(w_j|\Theta_c)^{x_i^j}$   
 $x_i^j$  décrit le nombre d'apparitions du mot  $j$  dans le document  $i$
- Notation:  $P(w_j|\Theta_c) \Rightarrow \Theta_c^j$ ,  
Résolution de:  $\Theta_c = \arg \max_{\Theta} \sum_{i=1}^{|C|} \sum_{j=1}^{|D|} x_i^j \log \Theta_c^j$
- Solution:  $\Theta_c^j = \frac{\sum_{d_i \in C} x_i^j}{\sum_{d_i \in C} \sum_{j \in D} x_i^j}$

## Naive Bayes (suite)

- Très simple à calculer (possibilité de travailler directement en base de données)
- Naturellement multi-classes,

$$\text{inférence: } \arg \max_c \sum_{j=1}^{|D|} x_i^j \log(\Theta_c^j)$$

Performance intéressante... Mais améliorable

Quid des mots inconnus dans une classe?

## Naive Bayes (suite)

- Très simple à calculer (possibilité de travailler directement en base de données)
- Naturellement multi-classes,

$$\text{inférence: } \arg \max_c \sum_{j=1}^{|D|} x_i^j \log(\Theta_c^j)$$

Performance intéressante... Mais améliorable

Quid des mots inconnus dans une classe?

- Extensions:

$$\text{Robustesse : } \Theta_c^j = \frac{\sum_{d_i \in C_m} x_i^j + \alpha}{\sum_{d_i \in C_m} \sum_{j \in D} x_i^j + \alpha |D|}$$

- Mots fréquents (*stopwords*) ... Bcp d'importance dans la décision
- pas d'aspect discriminant



## Classifieur linéaire, mode de décision

- Données en sacs de mots, différents codages possibles:

$$X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1d} \\ \vdots & & & \\ x_{N1} & x_{N2} & \cdots & x_{Nd} \end{bmatrix}$$

Un document  $\mathbf{x}_i \in \mathbb{R}^d$

# Classifieur linéaire, mode de décision

- Données en sacs de mots, différents codages possibles:

$$X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1d} \\ \vdots & & & \\ x_{N1} & x_{N2} & \cdots & x_{Nd} \end{bmatrix}$$

Un document  $\mathbf{x}_i \in \mathbb{R}^d$

- Décision linéaire:

- Décision linéaire simple:

$$f(\mathbf{x}_i) = \mathbf{x}_i \mathbf{w} = \sum_j x_{ij} w_j$$

- Régression logistique:

$$f(\mathbf{x}_i) = \frac{1}{1 + \exp(-(\mathbf{x}_i \mathbf{w} + b))}$$

# Classifieur linéaire, mode de décision

- Données en sacs de mots, différents codages possibles:

$$X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1d} \\ \vdots & & & \\ x_{N1} & x_{N2} & \cdots & x_{Nd} \end{bmatrix}$$

Un document  $\mathbf{x}_i \in \mathbb{R}^d$

- Décision linéaire:

- Décision linéaire simple:

$$f(\mathbf{x}_i) = \mathbf{x}_i \mathbf{w} = \sum_j x_{ij} w_j$$

- Régression logistique:

$$f(\mathbf{x}_i) = \frac{1}{1 + \exp(-(\mathbf{x}_i \mathbf{w} + b))}$$

- Mode de fonctionnement bi-classe (extension par un-contre-tous)

# Formulations et contraintes

## ■ Formulation

- Maximisation de la vraisemblance ( $y_i \in \{0, 1\}$ ):

$$L = \prod_{i=1}^N P(y_i = 1 | \mathbf{x}_i)^{y_i} \times [1 - P(y_i = 1 | \mathbf{x}_i)]^{1-y_i}$$

$$L_{\log} = \sum_{i=1}^N y_i \log(f(\mathbf{x}_i)) + (1 - y_i) \log(1 - f(\mathbf{x}_i))$$

- Minimisation d'un coût ( $y_i \in \{-1, 1\}$ ):

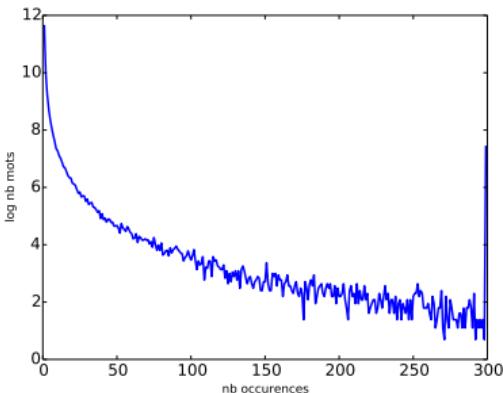
$$C = \sum_{i=1}^N (f(\mathbf{x}_i) - y_i)^2 \quad \text{ou} \quad C = \sum_{i=1}^N (-y_i f(\mathbf{x}_i))_+$$

- Passage à l'échelle: technique d'optimisation, gradient stochastique, calcul distribué
- Fléau de la dimensionnalité...

# Fléau de la dimensionnalité

$$X = \begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1d} \\ \vdots & & & \\ x_{N1} & x_{N2} & \cdots & x_{Nd} \end{bmatrix}$$

- $d = 10^5, N = 10^4\dots$
- Distribution des mots en fonction de leurs fréquences:



- Construire un système pour bien classer tous les documents proposés

## Fléau de la dimensionnalité (suite)

- Il est souvent (toujours) possible de trouver des mots qui n'apparaissent que dans l'une des classes de documents...
- Il suffit de se baser dessus pour prendre une décision parfaite...
- Mais ces mots apparaissent ils dans les documents non vus jusqu'ici???

# Régularisation

Idée:

Ajouter un terme sur la fonction coût (ou vraisemblance) pour pénaliser le nombre (ou le poids) des coefficients utilisés pour la décision

$$L_{\log} = \sum_{i=1}^N y_i \log(f(\mathbf{x}_i)) + (1 - y_i) \log(1 - f(\mathbf{x}_i)) - \lambda \|\mathbf{w}\|_\alpha$$

$$C = \sum_{i=1}^N (f(\mathbf{x}_i) - y_i)^2 + \lambda \|\mathbf{w}\|_\alpha, \quad C = \sum_{i=1}^N (-y_i f(\mathbf{x}_i))_+ + \lambda \|\mathbf{w}\|_\alpha$$

Avec:  $\|\mathbf{w}\|_2 = \sum_j w_j^2$  ou  $\|\mathbf{w}\|_1 = \sum_j |w_j|$

- Etude de la mise à jour dans un algorithme de gradient
- On se focalise sur les coefficients *vraiment* importants

# DONNÉES (EXEMPLES)

# Analyse de twitter (& réseaux sociaux en général)

- Sur un mot clés:
  - Quantification des documents positifs/négatifs
- Groupe politiques
- Emergence de thématiques...

# Locuteurs, Chirac/Mitterrand

Données d'apprentissage:

```
<100:1:C> Quand je dis chers amis, ...
<100:2:C> D'abord merci de cet ...
...
<100:14:M> Et ce sentiment ...
```

Le format est le suivant: {ID-Discours:ID-phrase:Etiquette}, C → Chirac, M → Mitterrand

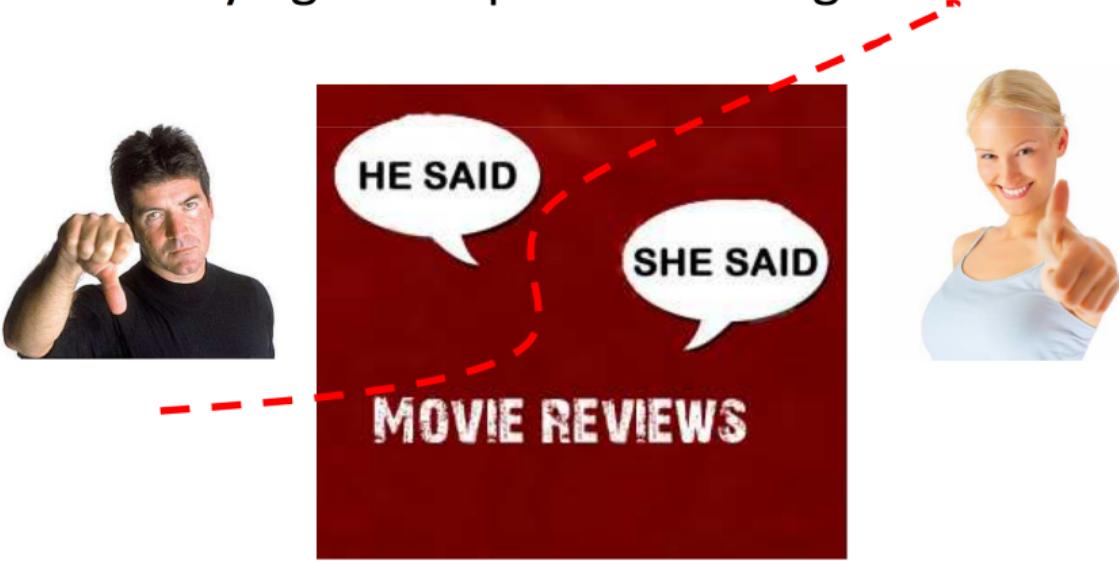
Données de test, sans les étiquettes:

```
<100:1> Quand je dis chers amis, ...
<100:2> D'abord merci de cet ...
...
```

# Revues de films

## The Task

Building a model for movies reviews in English for classifying it into positive or negative.



## Revues de films

# *Sentiment Polarity Dataset Version 2.0*

1000 positive movie review and 1000 negative review texts from:

***Thumbs up? Sentiment Classification using Machine Learning Techniques.*** Bo Pang, Lillian Lee, and Shivakumar Vaithyanathan. Proceedings of EMNLP, pp. 79--86, 2002.

“Our data **source** was the **Internet Movie Database** (IMDb) archive of the rec.arts.movies.reviews newsgroup.<sup>3</sup> We selected only reviews where the **author rating** was **expressed** either with stars or some **numerical value** (other conventions varied too widely to allow for automatic processing). Ratings were automatically extracted and converted into one of three categories: positive, negative, or neutral. For the work described in this paper, we concentrated **only** on discriminating between **positive** and **negative** sentiment.”

## Revues de films

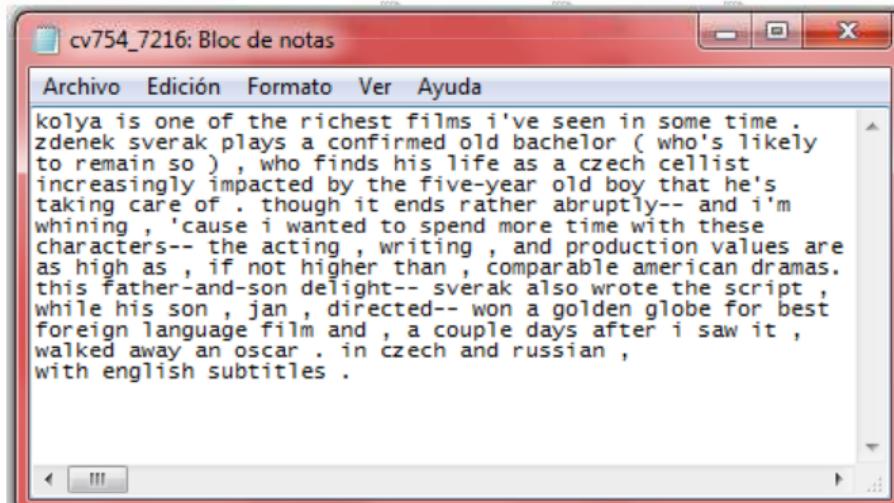
## The Data (1/2)



## Biblioteca Documentos

pos

cv754_7216	cv114_18398	cv938_10220	cv013_10159
cv280_8267	cv471_16858	cv424_8831	cv253_10077
cv825_5063	cv230_7428	cv763_14729	cv722_7110
cv057_7453	cv075_6500	cv319_14727	cv082_11080
cv640_5378	cv058_8025	cv430_17351	cv312_29377
			cv361_28944
			cv931_17563
			cv170_3006



# Compétition

UE TAL :

Obligation de participer à une mini-compétition sur les 2 jeux de données

⇒ Buts:

- Traiter des données textuelles (!)
- Travail minimum d'optimisation des classifieurs
- Post-traitements & interactions (minimales) avec un système externe

# EVALUATION/OUTILS

• Evaluation  
• Outils

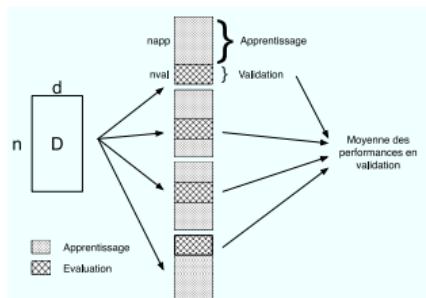
# Comment évaluer les performances?

## ■ Métriques d'évaluation

- Taux de reconnaissance  $\frac{N_{correct}}{N_{tot}}$
- Précision (dans la classe c)  $\frac{N_{correct}^c}{N_{predits}^c}$
- Rappel (dans la classe c) (=couverture)  $\frac{N_{correct}^c}{N_{tot}^c}$
- F1  $\frac{(1+\beta^2)precision \cdot rappel}{\beta^2 precision + rappel}$
- ROC (faux pos VS vrai pos) / AUC

## ■ Procédures

- Apprentissage/test



- Validation croisée
- Leave-one-out

# Analyse qualitative

Regarder les poids des mots du classifieur:

annoying	37.2593
another	-8.458
any	3.391
anyone	-1.4651
anything	-15.5326
anyway	29.2124
apparently	12.5416
...	
attention	-1.2901
audience	1.7331
audiences	-3.7323
away	-14.9303
awful	30.8509

# Ressources (python)

- nltk
  - Corpus, ressources, listes de *stopwords*
  - quelques classifieurs (mais moins intéressant que sklearn)
- gensim
  - Très bonne implémentation (rapide)
  - Outils pour la sémantique statistique (cours suivants)
- sklearn
  - Boite à outils de machine learning (SVM, Naive Bayes, regression logistique ...)
  - Evaluations diverses
  - Quelques outils pour le texte (simples mais pas très optimisés)

# Ressources systèmes

- Lancer des expériences à distance:
  - nohup (simple, mais perte du terminal)
    - Redirection, usage des logs
  - screen, tmux
- Connexion à distance = usage d'une passerelle
  - tunnel ssh
- Gestion des quotas
  - Travail sur le /tmp

⇒ Un rythme à trouver: fiabiliser le code en local, lancer les calculs lourds la nuit ou le week-end

# Aspects industriels

## 1 Récupération/importation d'un corpus

- Lecture de format XML
- Template NLTK...

## 2 Optimisation d'un modèle.

- Campagne d'expérience (d'abord grossière - codage, choix modèle...-, puis fine - régularisation...)
- Assez long... Mais essentielle
- **Le savoir-faire est ici**

## 3 Evaluation des performances (souvent en même temps que la phase d'optimisation)

- Usage de la validation croisée

## 4 Apprentissage + packaging du modèle final

- Définition des formats IO
- Mode de fonctionnement : API, service web...
- Documentation